

The Forward Slice Core Microarchitecture

Kartik Lakshminarasimhan (*Ghent University*)

Ajeya Naithani (*Ghent University*)

Josué Feliu Pérez (*University of Murcia*)


Lieven Eeckhout (*Ghent University*)



UNIVERSITEIT
GENT

October 7, 2020 -- PACT 2020

Forward Slice Core (FSC) Delivers High Performance at Low Complexity



Modern computing systems are energy-constrained

OoO cores deliver high performance but are complex

Goal: OoO performance at in-order complexity

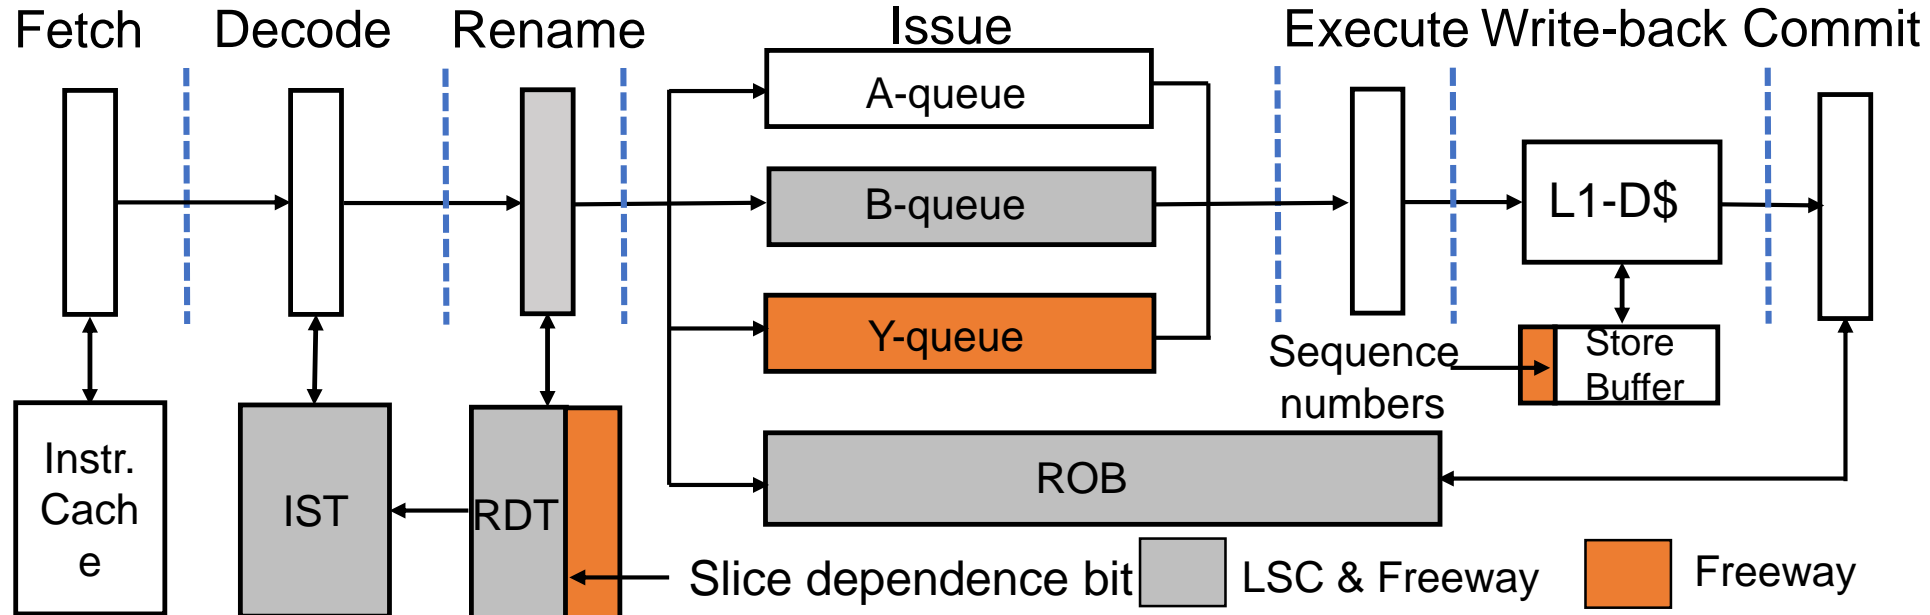
Comparison Against Prior Work

Processor	ILP	MHP	Complexity
InO	-	-	+
LSC <i>[ISCA '15]</i>	+	+	+
Freeway <i>[HPCA '19]</i>	+	++	+++
FSC	+++	+++	++
OoO	++++	++++	+++++

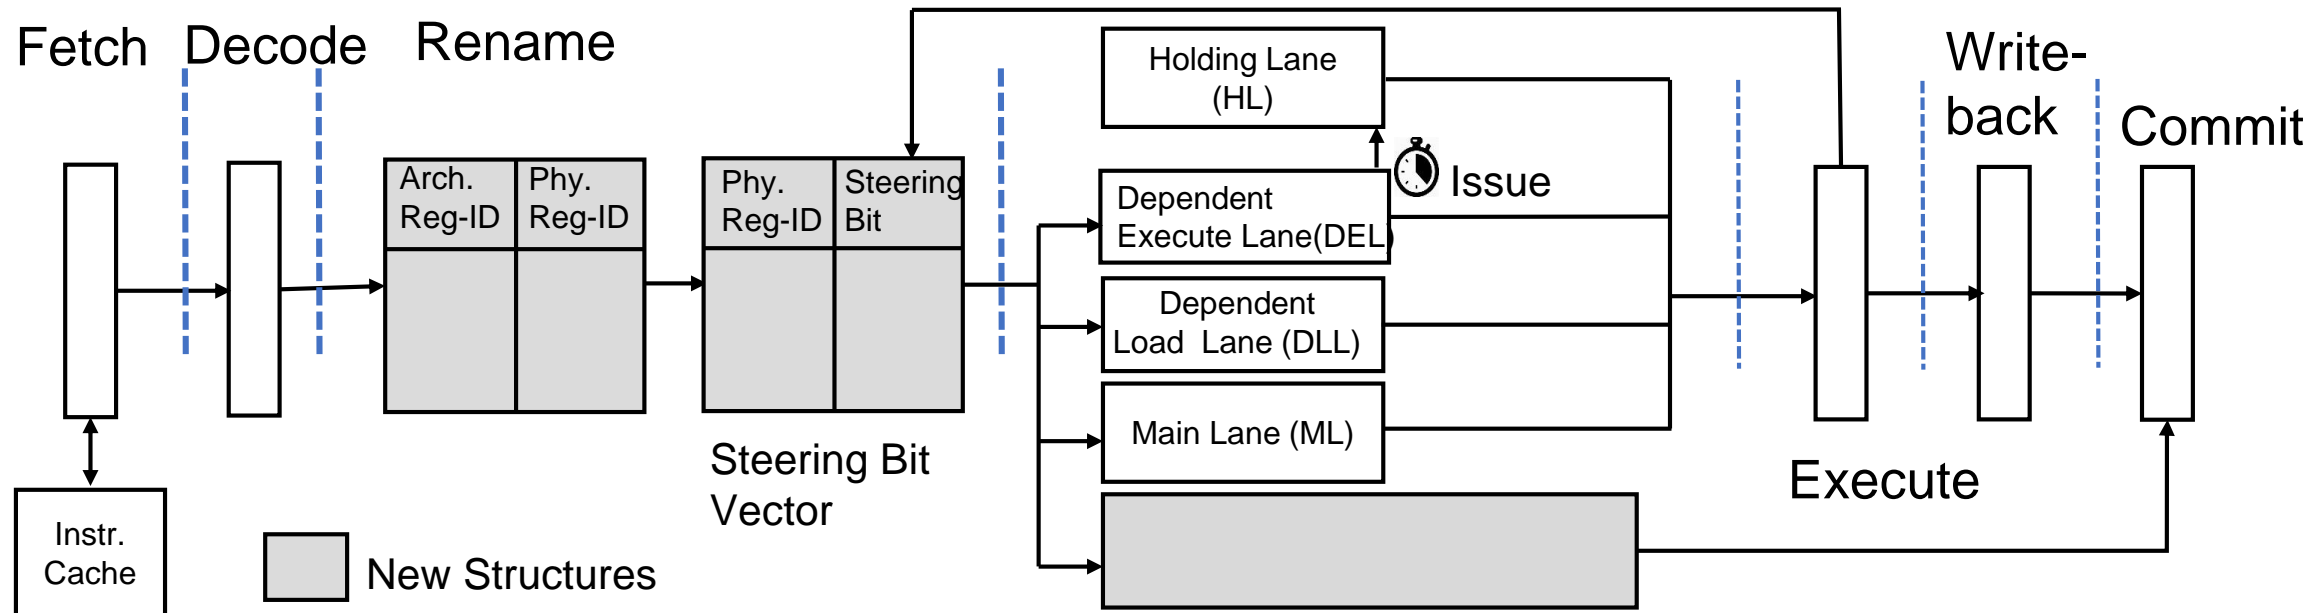
MHP = Memory Hierarchy Parallelism = number of overlapping memory accesses that hit anywhere in the memory hierarchy

Prior Work: Slice-Out-of-Order Processors

- Slice-out-of-order (sOoO) processors
 - Load Slice Core (LSC) [ISCA'15] and Freeway [HPCA'19]
- Limitations for sOoO processors:
 - Imperfect backward slice identification
 - Lack of ILP




The Forward Slice Core





FSC's key features:

- Restricted OoO execution through in-order lanes
- Identify and steer forward slice instructions
- Move stalling forward slice instructions to Holding Lane (HL)
- Replicate store address micro-ops across issue lanes for in-order memory address resolution



Forward Slice Example





a code example from mcf





Steering In-Flight Load Consumers to DEL and DLL Lanes





**load
consumers
are steered
to separate
lanes**



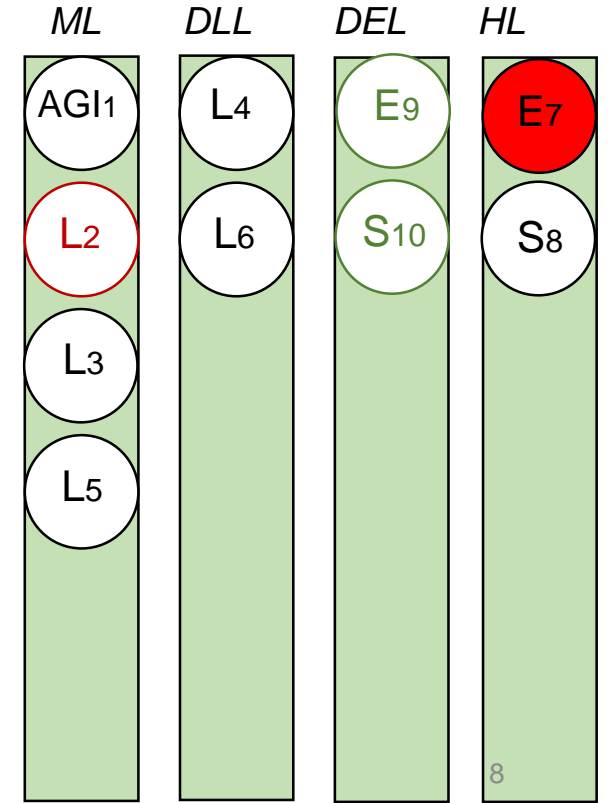

Redirecting Long-Latency Load Consumers to Holding Lane



load consumers are steered to separate lanes





re-directing L1D-miss load consumers from DEL lane to HL




Store Address Replication (SAR)

No memory dependence violations by
replicating store-address instructions
→ in-order address computation




Steering Bit Vector (SBV): Tracking In-Flight Load Consumers



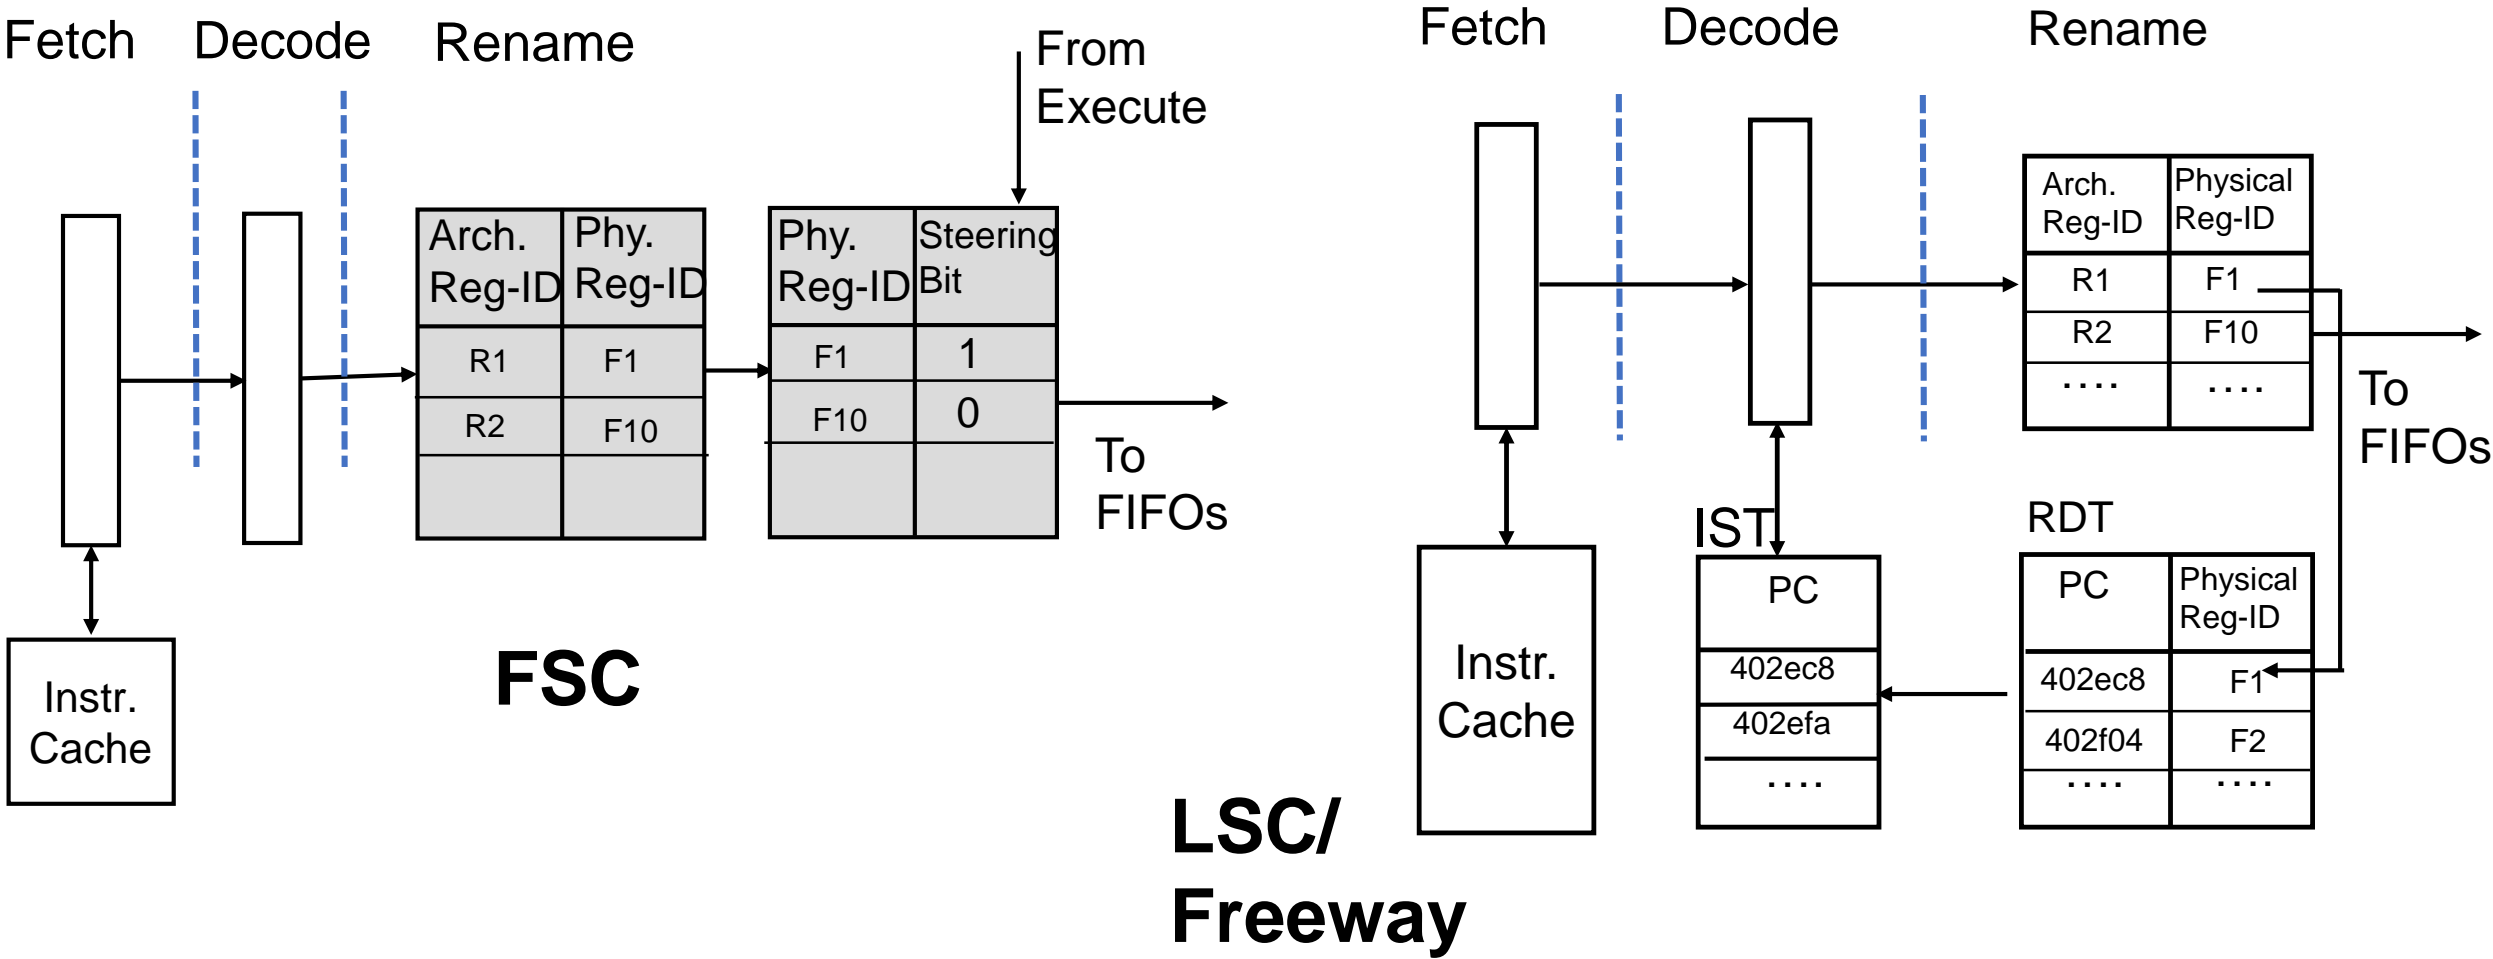
```
L1: ld ( r8 + r7*8 ), r5
L2: ld ( r5 + r6*8 ), r1
E3: add r1, r1
S4: st r1, (r2 + r3*4)
```

Steering Bit Vector (SBV)

Phy. Reg-ID	Steering Bit
F5	1 → 0
F1	1




Complexity of Steering Logic: FSC vs Freeway




FSC requires 1048 less bytes than Freeway and does not require CAM logic for memory disambiguation


Why Does FSC Outperform LSC and Freeway?




LSC



Freeway



FSC



Experimental Setup

Sniper multi-core simulator


- ARM Cortex-A7-like configuration
- 32KB L1, 512 KB L2
- 22 nm (McPAT and CACTI 6.5)



SPEC CPU2017 representative 1B instruction SimPoints

InO	LSC	Freeway	FSC	OoO
Stall-on-use	Restricted OoO	Restricted OoO	Restricted OoO	OoO
32-entry queue	16-entry A and B queues	12-entry A, B and Y queues	8-entry ML, DEL, DLL and HL	32-entry ROB and issue queue
2-wide issue	2-wide issue	2-wide issue	2-wide issue	2-wide issue
-	128-entry IST 64-entry RDT	128 entry IST 64-entry RDT	64-entry SBV	-

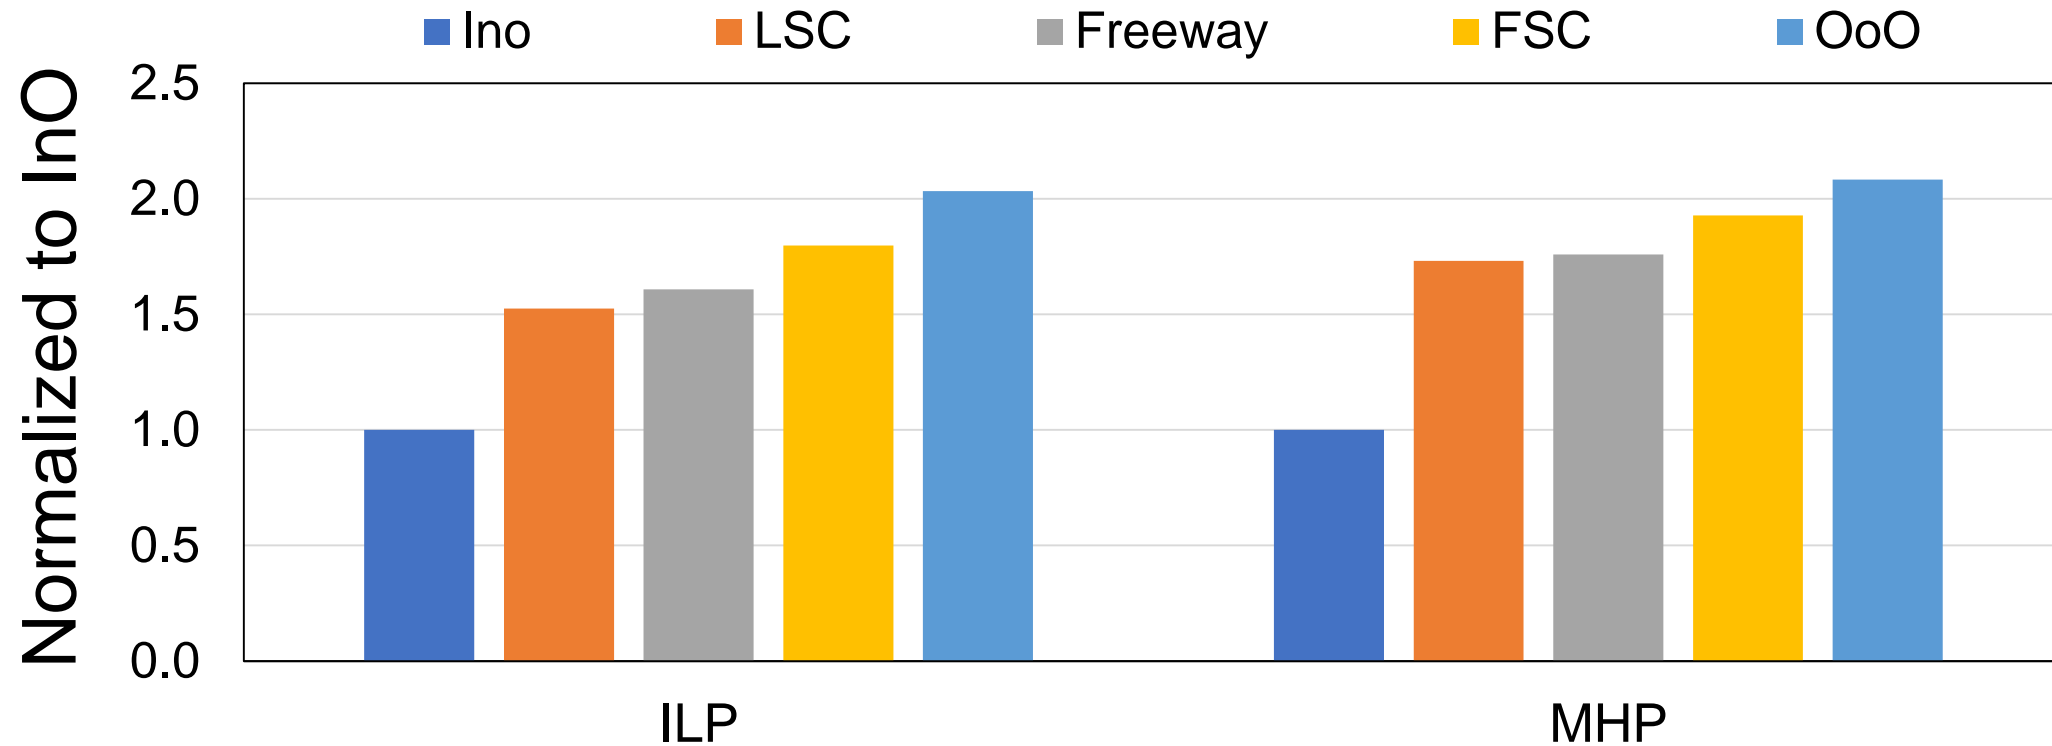
FSC Performance Benefit



9.7% avg performance improvement over Freeway with 1408 bytes less hardware


Within 6.9% of OoO, while consuming 56% less power

FSC Improves ILP and MHP



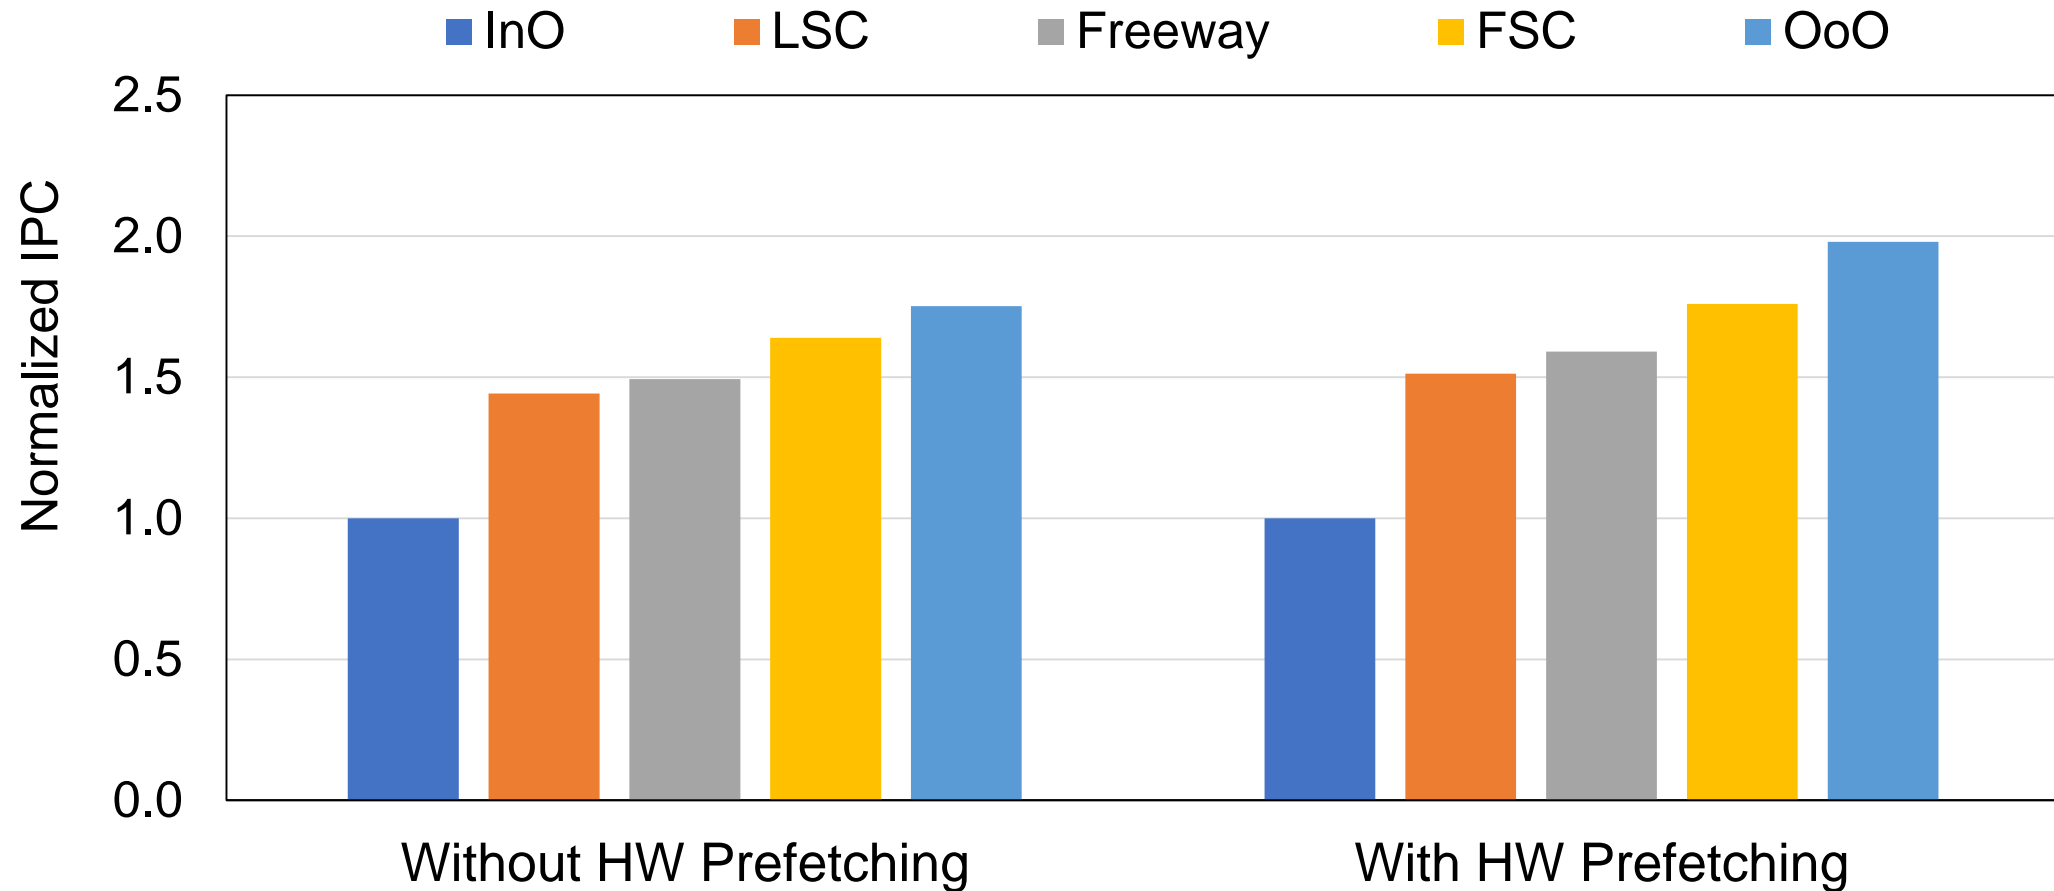
FSC improves ILP by 12% and MHP by 10% over Freeway

FSC Performance Benefit over CESP



FSC improves performance by 4.5% on average (and up to 12.6%) over CESP *[Palacharla et al. ISCA'97]* at reduced hardware complexity

FSC with Hardware Prefetcher



FSC improves performance by 75% on average compared to an InO core

Conclusion: Forward Slice Core

Three key innovations

- Steering load-consumers to separate in-order queues
- Redirect L1D-miss load-consumers to holding lane
- Store address replication

Higher ILP and MHP than prior sOoO cores, at reduced hardware cost and power consumption

Key results

- 9.7% higher performance than Freeway with less HW
- within 6.9% of OoO core with 56% less power consumption
- 44% higher performance than InO at small HW cost